Choose the representation for the device budget

MEMORY CHANGES THE RANKING

LUT vs Horner at 80 B: lower mean error
At 144 / 224 B: the ranking reverses

LUT SAVES COMPUTATION

Inference vs Horner: 13.2–13.8% less time
Projected updates vs cubic: 30.7–39.1% less

ACCURACY STILL COSTS

144 B: Horner improves MAE
0.05571 → 0.04552 °C

Synthetic calibration study • Physical Mega2560 / ESP32-C3 execution • Reference-based adaptation